

Machine Learning Homework 6

April 24, 2025

Gabe Morris

```
[1]: # toc
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from keras.api.models import Model
from keras.api.layers import Input, Dense
from keras.api.optimizers import Adam

plt.style.use('../maroon_ipynb.mplstyle')
tf.config.set_visible_devices([], 'GPU')
```

Contents

Problem 1	3
Solution	3
Generate Boundary and Collocation Points	3
Build the PINN Model	4
Training Step	5

Problem 1

A function $u(x, y)$ is defined on a unit square $x \in [0, 1]$, $y \in [0, 1]$, and obeys the following partial differential equation:

$$\nabla^2 u(x, y) = e^{-x} (x - a + y^3 + by)$$

Where a and b are constants. The function $u(x, y)$ is also subject to the following Dirichlet boundary conditions:

$$\begin{aligned} u(0, y) &= y^3 \\ u(1, y) &= (1 + y^3)/e \\ u(x, 0) &= xe^{-x} \\ u(x, 1) &= e^{-x}(1 + x) \end{aligned}$$

For particular values of a and b the analytic solution is:

$$u(x, y) = e^{-x}(x + y^3)$$

Construct a physics-informed neural network (PINN) to determine the unknown constants a and b which give this solution. Try to minimize the number of points away from the boundary which explicitly use the analytic solution. Plot how the prediction of a and b evolves with the number of training epochs.

Solution

To solve this problem, we need to determine the unknown constants a and b in a partial differential equation (PDE) using a Physics-Informed Neural Network (PINN). The solution involves training a neural network to approximate the function $u(x, y)$ while satisfying the given boundary conditions and the PDE, which includes the constants a and b . These constants are treated as trainable variables in the model.

Generate Boundary and Collocation Points

Boundary points are sampled from the edges of the unit square, and their corresponding u values are computed using the given Dirichlet boundary conditions. Collocation points are randomly sampled from the interior of the unit square to enforce the PDE constraint.

```
[2]: # Define some constant sizes
N_b = 100 # number of points per edge
N_f = 1000 # number of collocation points

# For x=0
x_b1 = np.zeros((N_b, 1))
y_b1 = np.linspace(0, 1, N_b).reshape(-1, 1)
u_b1 = y_b1**3
```

```

# For x=1
x_b2 = np.ones((N_b, 1))
y_b2 = np.linspace(0, 1, N_b).reshape(-1, 1)
u_b2 = (1 + y_b2**3)/np.exp(1)

# For y=0
x_b3 = np.linspace(0, 1, N_b).reshape(-1, 1)
y_b3 = np.zeros((N_b, 1))
u_b3 = x_b3*np.exp(-x_b3)

# For y=1
x_b4 = np.linspace(0, 1, N_b).reshape(-1, 1)
y_b4 = np.ones((N_b, 1))
u_b4 = np.exp(-x_b4)*(1 + x_b4)

# Combine all boundary points
X_bc = np.vstack([
    np.hstack([x_b1, y_b1]),
    np.hstack([x_b2, y_b2]),
    np.hstack([x_b3, y_b3]),
    np.hstack([x_b4, y_b4])
])
u_bc = np.vstack([u_b1, u_b2, u_b3, u_b4])

# Generate collocation points
x_f = np.random.rand(N_f, 1)
y_f = np.random.rand(N_f, 1)
X_pde = np.hstack([x_f, y_f])

# Convert to TensorFlow tensors
X_bc_tf = tf.convert_to_tensor(X_bc, dtype=tf.float32)
u_bc_tf = tf.convert_to_tensor(u_bc, dtype=tf.float32)
X_pde_tf = tf.convert_to_tensor(X_pde, dtype=tf.float32)

```

Build the PINN Model

```

[3]: def create_model():
    inputs = Input(shape=(2,))
    x = Dense(64, activation='tanh')(inputs)
    x = Dense(64, activation='tanh')(x)
    x = Dense(64, activation='tanh')(x)
    outputs = Dense(1)(x)
    model_ = Model(inputs=inputs, outputs=outputs)
    return model_

model = create_model()

```

```

# Initialize trainable variables a and b
a = tf.Variable(1.5, dtype=tf.float32, name='a')
b = tf.Variable(6.5, dtype=tf.float32, name='b')

# Define the optimizer
optimizer = Adam(learning_rate=0.001)

```

Training Step

The total loss (sum of boundary, PDE losses, and analytical/data loss) is minimized using the Adam optimizer. The gradients of the total loss with respect to the network's parameters and the constants a and b are computed and used to update these variables.

```

[4]: pde_loss_weight = 1 # weight to emphasize PDE constraint
analytic_loss_weight = 1 # weight to emphasize the analytic solution

@tf.function
def train_step(X_bc_, u_bc_, X_pde_):
    with tf.GradientTape() as tape:
        # BC loss
        u_pred_bc = model(X_bc_, training=True)
        bc_loss_ = tf.reduce_mean((u_bc_ - u_pred_bc)**2)

        # PDE loss
        with tf.GradientTape(persistent=True) as tape2:
            tape2.watch(X_pde_)
            u_pred_pde = model(X_pde_, training=True)
            grad_u = tape2.gradient(u_pred_pde, X_pde_)
            u_x = grad_u[:, 0:1]
            u_y = grad_u[:, 1:2]
            u_xx = tape2.gradient(u_x, X_pde_)[:, 0:1]
            u_yy = tape2.gradient(u_y, X_pde_)[:, 1:2]
        del tape2

        x_pde = X_pde_[:, 0:1]
        y_pde = X_pde_[:, 1:2]
        rhs = tf.exp(-x_pde)*(x_pde - a + y_pde**3 + b*y_pde)
        PDE_resid = u_xx + u_yy - rhs
        pde_loss_ = tf.reduce_mean(PDE_resid**2)

        # Analytic loss
        u_analytic = tf.exp(-x_pde)*(x_pde + y_pde**3)
        analytic_loss_ = tf.reduce_mean((u_analytic - u_pred_pde)**2)

        total_loss_ = bc_loss_ + pde_loss_weight*pde_loss_ +
        ↪ analytic_loss_weight*analytic_loss_

```

```

grads = tape.gradient(total_loss_, model.trainable_variables + [a, b])
optimizer.apply_gradients(zip(grads, model.trainable_variables + [a, b]))
return bc_loss_, pde_loss_, analytic_loss_, total_loss_

```

Now we can finish training and tracking the values of a and b .

```

[5]: epochs = 50_000
a_history, b_history = [], []
loss_history = []

for epoch in range(1, epochs + 1):
    # Resample interior
    x_f = np.random.rand(N_f, 1)
    y_f = np.random.rand(N_f, 1)
    X_pde_tf = tf.convert_to_tensor(np.hstack([x_f, y_f]), dtype=tf.float32)

    bc_loss, pde_loss, analytical_loss, total_loss = train_step(X_bc_tf,
↪u_bc_tf, X_pde_tf)
    a_history.append(a.numpy())
    b_history.append(b.numpy())

    if epoch%500 == 0:
        print(f"Epoch {epoch:5d} | BC_loss = {bc_loss:.3e} PDE_loss =
↪{pde_loss:.3e} A_loss = {analytical_loss:.3e} a = {a.numpy():.4f} b = {b.
↪numpy():.4f}")

epoch_array = np.arange(1, epochs + 1)
plt.plot(epoch_array, a_history, label='a')
plt.plot(epoch_array, b_history, label='b')
plt.axhline(y=2, ls='--', color='darkgrey')
plt.axhline(y=6, ls='--', color='darkgrey')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.legend()
plt.show()

```

WARNING:tensorflow:Calling GradientTape.gradient on a persistent tape inside its context is significantly less efficient than calling it outside the context (it causes the gradient ops to be recorded on the tape, leading to increased CPU and memory usage). Only call GradientTape.gradient inside the context if you actually want to trace the gradient in order to compute higher order derivatives.

```

Epoch   500 | BC_loss = 2.523e-03 PDE_loss = 2.926e-03 A_loss = 8.478e-04 a =
1.5470 b = 6.4472
Epoch  1000 | BC_loss = 1.960e-04 PDE_loss = 4.811e-04 A_loss = 4.051e-04 a =
1.5489 b = 6.4461
Epoch  1500 | BC_loss = 2.281e-04 PDE_loss = 1.315e-04 A_loss = 1.560e-04 a =

```

1.5520 b = 6.4442
 Epoch 2000 | BC_loss = 1.504e-04 PDE_loss = 7.817e-05 A_loss = 1.994e-04 a = 1.5562 b = 6.4416
 Epoch 2500 | BC_loss = 8.924e-05 PDE_loss = 1.404e-04 A_loss = 3.203e-04 a = 1.5616 b = 6.4382
 Epoch 3000 | BC_loss = 1.639e-04 PDE_loss = 2.183e-04 A_loss = 1.925e-04 a = 1.5686 b = 6.4337
 Epoch 3500 | BC_loss = 9.720e-05 PDE_loss = 5.953e-05 A_loss = 2.488e-04 a = 1.5772 b = 6.4280
 Epoch 4000 | BC_loss = 1.297e-04 PDE_loss = 1.516e-04 A_loss = 1.543e-04 a = 1.5878 b = 6.4210
 Epoch 4500 | BC_loss = 5.475e-04 PDE_loss = 1.046e-03 A_loss = 1.558e-04 a = 1.6006 b = 6.4123
 Epoch 5000 | BC_loss = 6.170e-05 PDE_loss = 1.109e-04 A_loss = 2.890e-04 a = 1.6154 b = 6.4021
 Epoch 5500 | BC_loss = 1.190e-04 PDE_loss = 1.494e-04 A_loss = 1.240e-04 a = 1.6321 b = 6.3900
 Epoch 6000 | BC_loss = 8.887e-05 PDE_loss = 1.317e-05 A_loss = 1.369e-04 a = 1.6506 b = 6.3762
 Epoch 6500 | BC_loss = 1.257e-04 PDE_loss = 2.180e-05 A_loss = 8.735e-05 a = 1.6702 b = 6.3614
 Epoch 7000 | BC_loss = 5.814e-05 PDE_loss = 1.132e-04 A_loss = 1.206e-04 a = 1.6904 b = 6.3454
 Epoch 7500 | BC_loss = 3.246e-04 PDE_loss = 1.850e-05 A_loss = 8.467e-05 a = 1.7107 b = 6.3290
 Epoch 8000 | BC_loss = 5.962e-05 PDE_loss = 1.141e-05 A_loss = 8.213e-05 a = 1.7306 b = 6.3125
 Epoch 8500 | BC_loss = 5.639e-05 PDE_loss = 7.340e-05 A_loss = 6.211e-05 a = 1.7500 b = 6.2963
 Epoch 9000 | BC_loss = 2.478e-05 PDE_loss = 8.141e-06 A_loss = 1.171e-04 a = 1.7684 b = 6.2806
 Epoch 9500 | BC_loss = 1.366e-04 PDE_loss = 3.675e-05 A_loss = 3.762e-05 a = 1.7849 b = 6.2665
 Epoch 10000 | BC_loss = 9.359e-05 PDE_loss = 2.256e-05 A_loss = 2.591e-05 a = 1.8016 b = 6.2521
 Epoch 10500 | BC_loss = 4.561e-05 PDE_loss = 4.385e-05 A_loss = 2.932e-05 a = 1.8174 b = 6.2386
 Epoch 11000 | BC_loss = 2.028e-05 PDE_loss = 2.000e-05 A_loss = 6.466e-05 a = 1.8326 b = 6.2252
 Epoch 11500 | BC_loss = 8.077e-05 PDE_loss = 1.023e-05 A_loss = 2.088e-05 a = 1.8461 b = 6.2138
 Epoch 12000 | BC_loss = 1.603e-05 PDE_loss = 5.899e-05 A_loss = 4.645e-05 a = 1.8592 b = 6.2026
 Epoch 12500 | BC_loss = 4.788e-05 PDE_loss = 9.831e-06 A_loss = 1.413e-05 a = 1.8713 b = 6.1919
 Epoch 13000 | BC_loss = 2.727e-05 PDE_loss = 4.291e-04 A_loss = 9.957e-05 a = 1.8830 b = 6.1818
 Epoch 13500 | BC_loss = 1.363e-05 PDE_loss = 7.390e-06 A_loss = 2.126e-05 a =

1.8941 b = 6.1718
 Epoch 14000 | BC_loss = 1.815e-05 PDE_loss = 4.107e-05 A_loss = 1.096e-05 a =
 1.9040 b = 6.1633
 Epoch 14500 | BC_loss = 1.698e-05 PDE_loss = 5.401e-06 A_loss = 8.622e-06 a =
 1.9132 b = 6.1552
 Epoch 15000 | BC_loss = 1.292e-05 PDE_loss = 4.562e-06 A_loss = 5.011e-05 a =
 1.9221 b = 6.1474
 Epoch 15500 | BC_loss = 7.260e-06 PDE_loss = 3.393e-06 A_loss = 1.083e-05 a =
 1.9301 b = 6.1401
 Epoch 16000 | BC_loss = 1.520e-04 PDE_loss = 2.280e-04 A_loss = 8.999e-05 a =
 1.9380 b = 6.1333
 Epoch 16500 | BC_loss = 2.419e-04 PDE_loss = 8.008e-04 A_loss = 3.620e-04 a =
 1.9452 b = 6.1268
 Epoch 17000 | BC_loss = 6.696e-05 PDE_loss = 9.631e-05 A_loss = 3.246e-05 a =
 1.9521 b = 6.1205
 Epoch 17500 | BC_loss = 1.578e-05 PDE_loss = 1.510e-05 A_loss = 3.904e-06 a =
 1.9586 b = 6.1144
 Epoch 18000 | BC_loss = 5.701e-06 PDE_loss = 3.523e-06 A_loss = 3.360e-06 a =
 1.9644 b = 6.1091
 Epoch 18500 | BC_loss = 2.878e-05 PDE_loss = 8.014e-05 A_loss = 8.526e-06 a =
 1.9698 b = 6.1040
 Epoch 19000 | BC_loss = 6.035e-06 PDE_loss = 4.117e-06 A_loss = 2.163e-06 a =
 1.9747 b = 6.0997
 Epoch 19500 | BC_loss = 2.787e-06 PDE_loss = 6.593e-05 A_loss = 6.186e-06 a =
 1.9794 b = 6.0952
 Epoch 20000 | BC_loss = 1.272e-04 PDE_loss = 3.409e-05 A_loss = 1.700e-04 a =
 1.9838 b = 6.0911
 Epoch 20500 | BC_loss = 2.037e-05 PDE_loss = 2.100e-04 A_loss = 3.687e-05 a =
 1.9877 b = 6.0873
 Epoch 21000 | BC_loss = 2.078e-06 PDE_loss = 8.640e-06 A_loss = 2.169e-06 a =
 1.9915 b = 6.0835
 Epoch 21500 | BC_loss = 1.174e-04 PDE_loss = 5.698e-04 A_loss = 1.518e-04 a =
 1.9948 b = 6.0803
 Epoch 22000 | BC_loss = 4.739e-05 PDE_loss = 5.263e-05 A_loss = 3.464e-05 a =
 1.9982 b = 6.0768
 Epoch 22500 | BC_loss = 2.504e-06 PDE_loss = 5.628e-06 A_loss = 4.995e-06 a =
 2.0012 b = 6.0738
 Epoch 23000 | BC_loss = 3.961e-04 PDE_loss = 3.128e-04 A_loss = 3.713e-04 a =
 2.0041 b = 6.0708
 Epoch 23500 | BC_loss = 1.909e-06 PDE_loss = 1.743e-06 A_loss = 5.357e-07 a =
 2.0066 b = 6.0682
 Epoch 24000 | BC_loss = 3.613e-06 PDE_loss = 6.637e-05 A_loss = 1.112e-06 a =
 2.0088 b = 6.0658
 Epoch 24500 | BC_loss = 2.808e-06 PDE_loss = 6.802e-06 A_loss = 4.889e-06 a =
 2.0109 b = 6.0635
 Epoch 25000 | BC_loss = 5.178e-06 PDE_loss = 1.301e-05 A_loss = 2.124e-06 a =
 2.0129 b = 6.0613
 Epoch 25500 | BC_loss = 7.016e-06 PDE_loss = 5.978e-05 A_loss = 6.861e-06 a =


2.0150 b = 6.0587
 Epoch 26000 | BC_loss = 7.190e-05 PDE_loss = 2.353e-05 A_loss = 8.153e-05 a =
 2.0166 b = 6.0569
 Epoch 26500 | BC_loss = 6.970e-07 PDE_loss = 1.213e-06 A_loss = 2.852e-07 a =
 2.0181 b = 6.0550
 Epoch 27000 | BC_loss = 3.775e-05 PDE_loss = 8.398e-05 A_loss = 3.638e-05 a =
 2.0195 b = 6.0533
 Epoch 27500 | BC_loss = 4.501e-06 PDE_loss = 2.283e-05 A_loss = 3.739e-06 a =
 2.0210 b = 6.0514
 Epoch 28000 | BC_loss = 1.483e-05 PDE_loss = 5.684e-06 A_loss = 1.680e-05 a =
 2.0219 b = 6.0503
 Epoch 28500 | BC_loss = 1.950e-06 PDE_loss = 1.262e-06 A_loss = 2.143e-06 a =
 2.0230 b = 6.0487
 Epoch 29000 | BC_loss = 1.154e-06 PDE_loss = 3.956e-06 A_loss = 6.364e-07 a =
 2.0239 b = 6.0475
 Epoch 29500 | BC_loss = 4.962e-07 PDE_loss = 1.571e-06 A_loss = 2.547e-07 a =
 2.0248 b = 6.0460
 Epoch 30000 | BC_loss = 2.103e-05 PDE_loss = 1.549e-04 A_loss = 1.681e-05 a =
 2.0256 b = 6.0448
 Epoch 30500 | BC_loss = 3.827e-06 PDE_loss = 1.435e-05 A_loss = 2.577e-06 a =
 2.0266 b = 6.0431
 Epoch 31000 | BC_loss = 4.923e-05 PDE_loss = 8.519e-06 A_loss = 4.839e-05 a =
 2.0272 b = 6.0422
 Epoch 31500 | BC_loss = 1.960e-04 PDE_loss = 6.677e-04 A_loss = 1.599e-04 a =
 2.0277 b = 6.0412
 Epoch 32000 | BC_loss = 1.201e-06 PDE_loss = 5.419e-06 A_loss = 6.793e-07 a =
 2.0282 b = 6.0402
 Epoch 32500 | BC_loss = 5.200e-07 PDE_loss = 1.414e-06 A_loss = 2.995e-07 a =
 2.0286 b = 6.0394
 Epoch 33000 | BC_loss = 4.013e-06 PDE_loss = 1.055e-04 A_loss = 1.650e-06 a =
 2.0290 b = 6.0385
 Epoch 33500 | BC_loss = 3.822e-06 PDE_loss = 1.401e-05 A_loss = 4.400e-06 a =
 2.0293 b = 6.0377
 Epoch 34000 | BC_loss = 9.144e-05 PDE_loss = 3.267e-03 A_loss = 8.340e-05 a =
 2.0298 b = 6.0366
 Epoch 34500 | BC_loss = 1.226e-06 PDE_loss = 5.008e-05 A_loss = 3.622e-07 a =
 2.0299 b = 6.0360
 Epoch 35000 | BC_loss = 3.813e-06 PDE_loss = 2.455e-06 A_loss = 3.015e-06 a =
 2.0301 b = 6.0352
 Epoch 35500 | BC_loss = 1.848e-05 PDE_loss = 3.194e-04 A_loss = 1.149e-05 a =
 2.0302 b = 6.0346
 Epoch 36000 | BC_loss = 1.354e-05 PDE_loss = 4.739e-06 A_loss = 1.267e-05 a =
 2.0304 b = 6.0338
 Epoch 36500 | BC_loss = 1.657e-06 PDE_loss = 4.506e-05 A_loss = 5.911e-07 a =
 2.0306 b = 6.0329
 Epoch 37000 | BC_loss = 2.053e-04 PDE_loss = 1.175e-05 A_loss = 1.949e-04 a =
 2.0306 b = 6.0324
 Epoch 37500 | BC_loss = 5.944e-05 PDE_loss = 2.986e-05 A_loss = 5.857e-05 a =

2.0307 b = 6.0317
 Epoch 38000 | BC_loss = 1.815e-06 PDE_loss = 4.910e-05 A_loss = 9.471e-07 a =
 2.0306 b = 6.0312
 Epoch 38500 | BC_loss = 1.659e-06 PDE_loss = 3.781e-06 A_loss = 1.659e-06 a =
 2.0306 b = 6.0306
 Epoch 39000 | BC_loss = 5.702e-06 PDE_loss = 1.697e-04 A_loss = 2.756e-06 a =
 2.0308 b = 6.0297
 Epoch 39500 | BC_loss = 2.551e-04 PDE_loss = 1.003e-04 A_loss = 2.700e-04 a =
 2.0307 b = 6.0292
 Epoch 40000 | BC_loss = 3.111e-07 PDE_loss = 3.129e-06 A_loss = 1.517e-07 a =
 2.0305 b = 6.0290
 Epoch 40500 | BC_loss = 1.205e-05 PDE_loss = 7.638e-06 A_loss = 1.077e-05 a =
 2.0305 b = 6.0283
 Epoch 41000 | BC_loss = 4.550e-05 PDE_loss = 3.737e-04 A_loss = 3.486e-05 a =
 2.0303 b = 6.0279
 Epoch 41500 | BC_loss = 2.956e-06 PDE_loss = 5.073e-06 A_loss = 2.116e-06 a =
 2.0302 b = 6.0274
 Epoch 42000 | BC_loss = 3.386e-07 PDE_loss = 2.169e-06 A_loss = 1.152e-07 a =
 2.0301 b = 6.0267
 Epoch 42500 | BC_loss = 1.151e-06 PDE_loss = 1.495e-06 A_loss = 5.762e-07 a =
 2.0301 b = 6.0263
 Epoch 43000 | BC_loss = 7.591e-06 PDE_loss = 2.480e-05 A_loss = 4.241e-06 a =
 2.0300 b = 6.0256
 Epoch 43500 | BC_loss = 5.173e-07 PDE_loss = 3.389e-06 A_loss = 7.160e-07 a =
 2.0297 b = 6.0254
 Epoch 44000 | BC_loss = 5.312e-07 PDE_loss = 2.657e-06 A_loss = 1.906e-07 a =
 2.0295 b = 6.0250
 Epoch 44500 | BC_loss = 1.101e-06 PDE_loss = 1.087e-06 A_loss = 5.777e-07 a =
 2.0294 b = 6.0245
 Epoch 45000 | BC_loss = 1.375e-05 PDE_loss = 2.747e-06 A_loss = 1.190e-05 a =
 2.0293 b = 6.0239
 Epoch 45500 | BC_loss = 3.930e-07 PDE_loss = 8.925e-07 A_loss = 5.347e-07 a =
 2.0290 b = 6.0237
 Epoch 46000 | BC_loss = 6.970e-06 PDE_loss = 4.998e-05 A_loss = 2.503e-06 a =
 2.0294 b = 6.0218
 Epoch 46500 | BC_loss = 2.383e-07 PDE_loss = 1.339e-06 A_loss = 7.805e-08 a =
 2.0288 b = 6.0223
 Epoch 47000 | BC_loss = 7.573e-07 PDE_loss = 1.431e-05 A_loss = 5.579e-07 a =
 2.0285 b = 6.0223
 Epoch 47500 | BC_loss = 1.762e-07 PDE_loss = 1.313e-06 A_loss = 1.387e-07 a =
 2.0282 b = 6.0220
 Epoch 48000 | BC_loss = 3.187e-05 PDE_loss = 5.253e-04 A_loss = 1.683e-05 a =
 2.0281 b = 6.0214
 Epoch 48500 | BC_loss = 1.834e-05 PDE_loss = 4.986e-06 A_loss = 1.618e-05 a =
 2.0279 b = 6.0211
 Epoch 49000 | BC_loss = 1.983e-07 PDE_loss = 1.061e-06 A_loss = 2.298e-07 a =
 2.0277 b = 6.0205
 Epoch 49500 | BC_loss = 2.709e-05 PDE_loss = 6.004e-06 A_loss = 2.443e-05 a =

2.0273 $b = 6.0205$

Epoch 50000 | BC_loss = $3.119\text{e-}06$ PDE_loss = $3.820\text{e-}05$ A_loss = $1.534\text{e-}06$ $a =$

2.0270 $b = 6.0202$



The above shows that the values for a and b converge to the expected values of 2 and 6. The analytical loss in the real world could be substituted for data loss.