

# Machine Learning Homework 6

April 24, 2025

Gabe Morris

```
[1]: # toc
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from keras.api.models import Model
from keras.api.layers import Input, Dense
from keras.api.optimizers import Adam

plt.style.use('../maroon_ipynb.mplstyle')
tf.config.set_visible_devices([], 'GPU')
```

## Contents

<b>Problem 1</b>	<b>3</b>
Solution . . . . .	3
Generate Boundary and Collocation Points . . . . .	3
Build the PINN Model . . . . .	4
Training Step . . . . .	5

## Problem 1

A function  $u(x, y)$  is defined on a unit square  $x \in [0, 1]$ ,  $y \in [0, 1]$ , and obeys the following partial differential equation:

$$\nabla^2 u(x, y) = e^{-x} (x - a + y^3 + by)$$

Where  $a$  and  $b$  are constants. The function  $u(x, y)$  is also subject to the following Dirichlet boundary conditions:

$$\begin{aligned} u(0, y) &= y^3 \\ u(1, y) &= (1 + y^3)/e \\ u(x, 0) &= xe^{-x} \\ u(x, 1) &= e^{-x}(1 + x) \end{aligned}$$

For particular values of  $a$  and  $b$  the analytic solution is:

$$u(x, y) = e^{-x}(x + y^3)$$

Construct a physics-informed neural network (PINN) to determine the unknown constants  $a$  and  $b$  which give this solution. Try to minimize the number of points away from the boundary which explicitly use the analytic solution. Plot how the prediction of  $a$  and  $b$  evolves with the number of training epochs.

## Solution

To solve this problem, we need to determine the unknown constants  $a$  and  $b$  in a partial differential equation (PDE) using a Physics-Informed Neural Network (PINN). The solution involves training a neural network to approximate the function  $u(x, y)$  while satisfying the given boundary conditions and the PDE, which includes the constants  $a$  and  $b$ . These constants are treated as trainable variables in the model.

### Generate Boundary and Collocation Points

Boundary points are sampled from the edges of the unit square, and their corresponding  $u$  values are computed using the given Dirichlet boundary conditions. Collocation points are randomly sampled from the interior of the unit square to enforce the PDE constraint.

```
[2]: # Define some constant sizes
N_b = 100 # number of points per edge
N_f = 1000 # number of collocation points

# For x=0
x_b1 = np.zeros((N_b, 1))
y_b1 = np.linspace(0, 1, N_b).reshape(-1, 1)
u_b1 = y_b1**3
```

```

# For x=1
x_b2 = np.ones((N_b, 1))
y_b2 = np.linspace(0, 1, N_b).reshape(-1, 1)
u_b2 = (1 + y_b2**3)/np.exp(1)

# For y=0
x_b3 = np.linspace(0, 1, N_b).reshape(-1, 1)
y_b3 = np.zeros((N_b, 1))
u_b3 = x_b3*np.exp(-x_b3)

# For y=1
x_b4 = np.linspace(0, 1, N_b).reshape(-1, 1)
y_b4 = np.ones((N_b, 1))
u_b4 = np.exp(-x_b4)*(1 + x_b4)

# Combine all boundary points
X_bc = np.vstack([
    np.hstack([x_b1, y_b1]),
    np.hstack([x_b2, y_b2]),
    np.hstack([x_b3, y_b3]),
    np.hstack([x_b4, y_b4])
])
u_bc = np.vstack([u_b1, u_b2, u_b3, u_b4])

# Generate collocation points
x_f = np.random.rand(N_f, 1)
y_f = np.random.rand(N_f, 1)
X_pde = np.hstack([x_f, y_f])

# Convert to TensorFlow tensors
X_bc_tf = tf.convert_to_tensor(X_bc, dtype=tf.float32)
u_bc_tf = tf.convert_to_tensor(u_bc, dtype=tf.float32)
X_pde_tf = tf.convert_to_tensor(X_pde, dtype=tf.float32)

```

## Build the PINN Model

```

[3]: def create_model():
    inputs = Input(shape=(2,))
    x = Dense(64, activation='tanh')(inputs)
    x = Dense(64, activation='tanh')(x)
    x = Dense(64, activation='tanh')(x)
    outputs = Dense(1)(x)
    model_ = Model(inputs=inputs, outputs=outputs)
    return model_

model = create_model()

```

```

# Initialize trainable variables a and b
a = tf.Variable(1.5, dtype=tf.float32, name='a')
b = tf.Variable(6.5, dtype=tf.float32, name='b')

# Define the optimizer
optimizer = Adam(learning_rate=0.005)

```

## Training Step

The total loss (sum of boundary and PDE losses) is minimized using the Adam optimizer. The gradients of the total loss with respect to the network's parameters and the constants  $a$  and  $b$  are computed and used to update these variables.

```

[4]: pde_loss_weight = 1 # weight to emphasize PDE constraint

@tf.function
def train_step(X_bc_, u_bc_, X_pde_):
    with tf.GradientTape() as tape:
        # BC loss
        u_pred_bc = model(X_bc_, training=True)
        bc_loss_ = tf.reduce_mean((u_bc_ - u_pred_bc)**2)

        # PDE loss (as before) ...
        with tf.GradientTape(persistent=True) as tape2:
            tape2.watch(X_pde_)
            u_pred_pde = model(X_pde_, training=True)
            grad_u = tape2.gradient(u_pred_pde, X_pde_)
            u_x = grad_u[:, 0:1]
            u_y = grad_u[:, 1:2]
            u_xx = tape2.gradient(u_x, X_pde_)[:, 0:1]
            u_yy = tape2.gradient(u_y, X_pde_)[:, 1:2]
        del tape2

        x_pde = X_pde_[:, 0:1]
        y_pde = X_pde_[:, 1:2]
        rhs = tf.exp(-x_pde)*(x_pde - a + y_pde**3 + b*y_pde)
        PDE_resid = u_xx + u_yy - rhs
        pde_loss_ = tf.reduce_mean(PDE_resid**2)

        total_loss_ = bc_loss_ + pde_loss_weight*pde_loss_

    grads = tape.gradient(total_loss_, model.trainable_variables + [a, b])
    optimizer.apply_gradients(zip(grads, model.trainable_variables + [a, b]))
    return bc_loss_, pde_loss_, total_loss_

```

Now we can finish training and tracking the values of  $a$  and  $b$ .

```

[5]: epochs = 10_000
a_history, b_history = [], []
loss_history = []

for epoch in range(1, epochs + 1):
    # Resample interior
    x_f = np.random.rand(N_f, 1)
    y_f = np.random.rand(N_f, 1)
    X_pde_tf = tf.convert_to_tensor(np.hstack([x_f, y_f]), dtype=tf.float32)

    bc_loss, pde_loss, total_loss = train_step(X_bc_tf, u_bc_tf, X_pde_tf)
    a_history.append(a.numpy())
    b_history.append(b.numpy())

    if epoch%500 == 0:
        print(f"Epoch {epoch:5d} | BC_loss = {bc_loss:.3e} PDE_loss = {pde_loss:.3e} a = {a.numpy():.4f} b = {b.numpy():.4f}")

epoch_array = np.arange(1, epochs + 1)
plt.plot(epoch_array, a_history, label='a')
plt.plot(epoch_array, b_history, label='b')
plt.axhline(y=2, ls='--')
plt.axhline(y=6, ls='--')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.legend()
plt.show()

```

WARNING:tensorflow:Calling GradientTape.gradient on a persistent tape inside its context is significantly less efficient than calling it outside the context (it causes the gradient ops to be recorded on the tape, leading to increased CPU and memory usage). Only call GradientTape.gradient inside the context if you actually want to trace the gradient in order to compute higher order derivatives.

Epoch	500		BC_loss = 4.528e-04	PDE_loss = 4.010e-04	a = 1.5998	b = 6.3810
Epoch	1000		BC_loss = 1.602e-04	PDE_loss = 9.751e-05	a = 1.6051	b = 6.3740
Epoch	1500		BC_loss = 9.363e-05	PDE_loss = 5.798e-05	a = 1.6119	b = 6.3644
Epoch	2000		BC_loss = 1.209e-04	PDE_loss = 3.179e-05	a = 1.6205	b = 6.3500
Epoch	2500		BC_loss = 4.884e-04	PDE_loss = 4.162e-04	a = 1.6275	b = 6.3390
Epoch	3000		BC_loss = 1.670e-04	PDE_loss = 1.066e-03	a = 1.6357	b = 6.3259
Epoch	3500		BC_loss = 1.717e-04	PDE_loss = 2.056e-04	a = 1.6450	b = 6.3097
Epoch	4000		BC_loss = 1.269e-04	PDE_loss = 6.969e-05	a = 1.6529	b = 6.2923
Epoch	4500		BC_loss = 5.008e-04	PDE_loss = 4.085e-04	a = 1.6631	b = 6.2711
Epoch	5000		BC_loss = 1.478e-04	PDE_loss = 1.140e-03	a = 1.6666	b = 6.2662
Epoch	5500		BC_loss = 4.289e-05	PDE_loss = 1.591e-05	a = 1.6732	b = 6.2501
Epoch	6000		BC_loss = 2.866e-04	PDE_loss = 6.171e-04	a = 1.6794	b = 6.2373
Epoch	6500		BC_loss = 6.570e-05	PDE_loss = 2.586e-05	a = 1.6857	b = 6.2173
Epoch	7000		BC_loss = 2.470e-04	PDE_loss = 3.796e-03	a = 1.6912	b = 6.2072

Epoch	7500		BC_loss = 8.221e-05	PDE_loss = 5.818e-05	a = 1.6958	b = 6.1933
Epoch	8000		BC_loss = 4.678e-05	PDE_loss = 7.445e-05	a = 1.7002	b = 6.1813
Epoch	8500		BC_loss = 2.385e-05	PDE_loss = 3.902e-05	a = 1.7042	b = 6.1700
Epoch	9000		BC_loss = 5.288e-05	PDE_loss = 6.377e-04	a = 1.7080	b = 6.1570
Epoch	9500		BC_loss = 4.551e-05	PDE_loss = 8.667e-05	a = 1.7115	b = 6.1447
Epoch	10000		BC_loss = 1.138e-04	PDE_loss = 2.260e-04	a = 1.7151	b = 6.1324

